



AI Cluster Architecture (SLURM + Kubernetes + RDMA + Parallel FS)

Use this as a quick, whiteboard-ready guide you can sketch in 60–120 seconds during the interview. Keep it simple, label clearly, and narrate the data/control flows.

1) High-Level Diagram (ASCII you can replicate fast)





Legend: solid arrows = control/API; bold blocks = data paths; RDMA fabric for NCCL collectives; Mgmt Ethernet for SLURM/K8s control.

2) 60-Second Walkthrough Script

- **Users → Head Nodes:** “Researchers submit via SLURM or K8s portals on HA head nodes running `slurmctld/slurmdbd` and K8s control plane. IAM handles SSO; Vault manages secrets.”
 - **Scheduling:** “SLURM owns GPUs using `select/cons_tres` and GRES; K8s is used for pipelines/operators. We integrate via SLURM Operator or run K8s workloads inside SLURM allocations.”
 - **Data:** “Cold/curated data in S3; hot POSIX via BeeGFS/GPFS/Lustre. Jobs optionally stage to node-local NVMe scratch.”
 - **Network:** “NDR/HDR InfiniBand spine–leaf for NCCL; GPU Direct RDMA enabled; topology-aware scheduling keeps jobs within a pod/leaf.”
 - **Observability:** “We track GPU/IB/FS KPIs with Prometheus + Grafana; SLURM accounting informs fair-share and capacity planning.”
-

3) SLURM Config Highlights (sound bites)

- `SelectType=select/cons_tres, GresTypes=gpu` with per-node `gres.conf` entries (e.g., `Name=gpu Type=mi300 File=/dev/dri/renderD128 Count=8`).
 - Partitions per workload (`ai-train, ai-infer, hpc`), QoS with preemption for urgent jobs.
 - Topology plugin + switches to prefer single-leaf placement; `--gpu-bind=closest`.
 - Accounting: `slurmdbd` + MySQL/Postgres; enforce usage caps and fair-share multifactor.
 - Container: `Prolog/Epilog` to mount datasets, set NCCL env, and collect job metrics.
-

4) Kubernetes Integration Patterns

Pattern A (K8s-primary): SLURM Operator/Volcano; K8s schedules workflow DAGs, offloads GPU allocations to SLURM.

Pattern B (SLURM-primary): Launch ephemeral K8s namespaces inside SLURM jobs for ML pipelines; preserves HPC env, good for multi-tenant clusters.

GPU mgmt: Expose GPUs to K8s via device plugins; align with SLURM GRES to avoid overcommit. Prefer Apptainer for MPI/HPC compatibility.

5) Storage Layout & Data Flow

- **Tiers:** S3/object (cold) → Parallel FS (warm) → Node-local NVMe (hot scratch).
- **Throughput:** Stripe large training data across OSTs/targets (e.g., 8–16); enable metadata striping (Lustre) for small-file directories.
- **Data lifecycle:** Pre-stage with Prolog; purge scratch on Epilog; use ILM (GPFS) or policies (BeeGFS) for tiering.

Training Flow: Loader pulls from POSIX → optional cache on NVMe → GPUs compute → checkpoints to POSIX/S3.

6) RDMA & NCCL Tuning Pointers

- IB NDR/HDR, PFC/ECN tuned; enable GPU Direct RDMA.
 - NCCL env: NCCL_IB_HCA, NCCL_IB_GID_INDEX, NCCL_NET_GDR_LEVEL, NCCL_TOPO_FILE for custom topo.
 - Pin ranks to GPUs/NUMA; ensure hwloc awareness.
-

7) Security & Multi-Tenancy

- SSO via OIDC; per-tenant accounts, partitions, and storage quotas.
 - Network isolation (VRFs/VLANs) between mgmt and data fabrics.
 - Image attestation for containers; secrets via Vault; read-only dataset mounts.
-

8) Observability & SLOs (what to say if asked)

- **GPU:** utilization, memory, SM occupancy.
- **Network:** IB port errors, retransmits, fabric load.
- **Storage:** metadata ops/s, read/write BW, client cache hits.

- **Scheduler:** queue wait vs run time, backfill efficiency.

Prometheus exporters (SLURM, node, IB, FS) + Grafana dashboards; alerts on underutilization and tail latencies.

9) Capacity & Cost Cheatsheet

- Rule of thumb: plan POSIX BW at ~2–3 GB/s per training node for large models; metadata IOPS sized for loaders.
 - At least 1:1 NVMe scratch to active dataset working set per node.
 - IB oversubscription \leq 2:1 for heavy All-Reduce; keep training jobs leaf-local when possible.
-

10) Example Command Snippets (verbal anchors)

- Submit 8-GPU job pinned within one node:

```
srun -N1 --gres=gpu:mi300:8 --gpu-bind=closest --ntasks-per-node=8
python train.py
```
 - Multi-node PyTorch:

```
srun -N8 --gres=gpu:mi300:8 --cpus-per-task=8 bash -c 'torchrun --
nproc_per_node=8 --nnodes=8 --rdzv_backend=c10d --rdzv_id=$SLURM_JOB_ID
--rdzv_endpoint=$SLURM_NODELIST train.py'
```
 - Pre-stage dataset (Prolog):

```
rsync /fs/datasets/$JOB_USER/$JOB_ID /local_nvme/datasets
```
-

11) Common Pitfalls to Call Out

- Fragmentation from overly specific constraints (features/GPUs) → fix with flexible requests and backfill.
 - Storage metadata hot spots → directory sharding, metadata striping.
 - K8s and SLURM double-scheduling GPUs → single source of truth for allocation.
 - NCCL hanging due to mixed GID or miswired topo → validate with `nccl-tests`.
-

12) 30-Second Closing

“Design keeps SLURM as the GPU source of truth, couples RDMA-aware placement with hot POSIX and node-local scratch, and integrates Kubernetes where it shines: pipelines and

automation. Observability and fair-share policies ensure high utilization without sacrificing researcher velocity.”